

# Peer Robotics Interview Assignment

## Objective:

This assignment tests your computer vision, ROS and programming skills. Develop a pallet detection & ground segmentation application in ROS2 for a manufacturing or warehousing environment. The solution should be optimized for deployment on edge devices like the NVIDIA Jetson AGX Orin, ensuring real-time performance suitable for mobile robotics applications.

## Tasks:

### 1. Dataset Acquisition and Preparation:

- **Dataset Acquisition:** Use the link to access open source database for pallets in different scenarios { [Pallets](#) }. You can find and use other open datasets for manufacturing/warehouse environments to further generalize your model.
- **Data Preparation:**
  - Annotate **pallets** & **ground** using existing annotation tools (ex. DINO)
  - Organize the dataset into training, validation, and test sets.
  - Apply data augmentation techniques (e.g., varying lighting conditions) to simulate real-world scenarios.

### 2. Object Detection and Semantic Segmentation:

- **Model Development:**
  - Train an object detection model (e.g., **YOLOv11** etc.) to detect pallets (GMA & Euro Pallets).
  - Train a semantic segmentation model to segment ground
  - Report training and validation stats.
- **Performance Evaluation:**
  - Assess the models using metrics like mAP for detection and IoU.
  - Ensure models perform robustly under varying environmental conditions.

### 3. ROS2 Node Development:

- **ROS2:**
  - Develop a ROS2 Python Node to run inference on the model.
  - Nodes should:
    - Subscribe to image and depth topics from a simulated or real camera.
    - Run inference on the input image for pallet detection & ground segmentation, and publish detections/segmentation overlaid on the input image.

#### 4. Edge Deployment Optimization (Optional) :

- **Model Optimization:**
  - Convert your models to formats suitable for edge deployment (e.g., **TensorRT**).
  - Apply optimization techniques like quantization and pruning to enhance performance.
- **Docker Container:**
  - Dockerized the complete module that can natively run on different devices as long as Nvidia drivers are present.

#### 5. Evaluation Criteria:

- We will run the module on AGX Orin and feed live camera data from zed 2i
- Performance will be measured based on pallet detection accuracy under varying conditions

**To streamline the assessment create a git repository with:**

1. README
2. Requirements.txt for all dependencies
3. Python file/notebook used for training
4. Link to the trained model
5. ros2 humble inference node which works off the shelf with the following [camera data](#) with output data overlaid on the original image:
  - a. Pallet Detection published on /pallet\_detection
  - b. Ground Segmentation published on /ground\_segmentation
6. [Optional] Python script that optimizes your model for a given GPU